

Mempersiapkan

```
# For this notebook to run with updated APIs, we need torch 1.12+ and torchvision 0.13+
```

```
try:
    import torch
    import torchvision
    assert int(torch.__version__.split(".")[1]) >= 12 or
int(torch.__version__.split(".")[0]) == 2, "torch version should be 1.12+"
    assert int(torchvision.__version__.split(".")[1]) >= 13,
"torchvision version should be 0.13+"
    print(f"torch version: {torch.__version__}")
    print(f"torchvision version: {torchvision.__version__}")
except:
    print(f"[INFO] torch/torchvision versions not as required,
installing nightly versions.")
    !pip3 install -U torch torchvision torchaudio --index-url
https://download.pytorch.org/whl/cu118
    import torch
    import torchvision
    print(f"torch version: {torch.__version__}")
    print(f"torchvision version: {torchvision.__version__}")
```

```
torch version: 2.1.0+cu118
torchvision version: 0.16.0+cu118
```

```
# Continue with regular imports
```

```
import matplotlib.pyplot as plt
import torch
import torchvision
```

```
from torch import nn
from torchvision import transforms
```

```
# Try to get torchinfo, install it if it doesn't work
```

```
try:
    from torchinfo import summary
except:
    print("[INFO] Couldn't find torchinfo... installing it.")
    !pip install -q torchinfo
    from torchinfo import summary
```

```
# Try to import the going_modular directory, download it from GitHub
if it doesn't work
```

```
try:
    from going_modular.going_modular import data_setup, engine
    from helper_functions import download_data, set_seeds,
plot_loss_curves
except:
```

```

# Get the going_modular scripts
print("[INFO] Couldn't find going_modular or helper_functions
scripts... downloading them from GitHub.")
!git clone https://github.com/mrdbourke/pytorch-deep-learning
!mv pytorch-deep-learning/going_modular .
!mv pytorch-deep-learning/helper_functions.py . # get the
helper_functions.py script
!rm -rf pytorch-deep-learning
from going_modular.going_modular import data_setup, engine
from helper_functions import download_data, set_seeds,
plot_loss_curves

[INFO] Couldn't find torchinfo... installing it.
[INFO] Couldn't find going_modular or helper_functions scripts...
downloading them from GitHub.
Cloning into 'pytorch-deep-learning'...
remote: Enumerating objects: 4033, done.ote: Counting objects: 100%
(1224/1224), done.ote: Compressing objects: 100% (225/225), done.ote:
Total 4033 (delta 1067), reused 1097 (delta 996), pack-reused 2809

device = "cuda" if torch.cuda.is_available() else "cpu"
device

'cuda'

```

Dapatkan Data

```

# Download pizza, steak, sushi images from GitHub
image_path =
download_data(source="https://github.com/mrdbourke/pytorch-deep-
learning/raw/main/data/pizza_steak_sushi.zip",
              destination="pizza_steak_sushi")
image_path

[INFO] data/pizza_steak_sushi directory exists, skipping download.

PosixPath('data/pizza_steak_sushi')

# Setup directory paths to train and test images
train_dir = image_path / "train"
test_dir = image_path / "test"

```

Buat Kumpulan Data dan Pemuat Data

Siapkan transformasi untuk gambar

```

# Create image size (from Table 3 in the ViT paper)
IMG_SIZE = 224

# Create transform pipeline manually

```

```

manual_transforms = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ToTensor(),
])
print(f"Manually created transforms: {manual_transforms}")

Manually created transforms: Compose(
  Resize(size=(224, 224), interpolation=bilinear, max_size=None,
  antialias=None)
  ToTensor()
)

```

Ubah gambar menjadi DataLoader

```

# Set the batch size
BATCH_SIZE = 32 # this is lower than the ViT paper but it's because
we're starting small

# Create data loaders
train_dataloader, test_dataloader, class_names =
data_setup.create_data_loaders(
    train_dir=train_dir,
    test_dir=test_dir,
    transform=manual_transforms, # use manually created transforms
    batch_size=BATCH_SIZE
)

train_dataloader, test_dataloader, class_names

(<torch.utils.data.dataloader.DataLoader at 0x7f18845ff0d0>,
<torch.utils.data.dataloader.DataLoader at 0x7f17f3f5f520>,
['pizza', 'steak', 'sushi'])

```

Visualisasikan satu gambar

```

# Get a batch of images
image_batch, label_batch = next(iter(train_dataloader))

# Get a single image from the batch
image, label = image_batch[0], label_batch[0]

# View the batch shapes
image.shape, label

(torch.Size([3, 224, 224]), tensor(2))

# Plot image with matplotlib
plt.imshow(image.permute(1, 2, 0)) # rearrange image dimensions to
suit matplotlib [color_channels, height, width] -> [height, width,
color_channels]

```

```
plt.title(class_names[label])  
plt.axis(False);
```



Mereplikasi makalah ViT: gambaran umum

Input dan output, lapisan dan blok

Lebih spesifik: ViT terbuat dari apa?

Menghitung bentuk input dan output yang menyematkan patch dengan tangan

```
# Create example values  
height = 224 # H ("The training resolution is 224.")  
width = 224 # W  
color_channels = 3 # C  
patch_size = 16 # P  
  
# Calculate N (number of patches)  
number_of_patches = int((height * width) / patch_size**2)  
print(f"Number of patches (N) with image height (H={height}), width  
(W={width}) and patch size (P={patch_size}): {number_of_patches}")  
  
Number of patches (N) with image height (H=224), width (W=224) and  
patch size (P=16): 196  
  
# Input shape (this is the size of a single image)  
embedding_layer_input_shape = (height, width, color_channels)
```

```
# Output shape
embedding_layer_output_shape = (number_of_patches, patch_size**2 *
color_channels)

print(f"Input shape (single 2D image): {embedding_layer_input_shape}")
print(f"Output shape (single 2D image flattened into patches):
{embedding_layer_output_shape}")

Input shape (single 2D image): (224, 224, 3)
Output shape (single 2D image flattened into patches): (196, 768)
```


Mengubah satu gambar menjadi tambalan

```
# View single image
plt.imshow(image.permute(1, 2, 0)) # adjust for matplotlib
plt.title(class_names[label])
plt.axis(False);
```



```
# Change image shape to be compatible with matplotlib (color_channels,
height, width) -> (height, width, color_channels)
image_permuted = image.permute(1, 2, 0)

# Index to plot the top row of patched pixels
patch_size = 16
plt.figure(figsize=(patch_size, patch_size))
plt.imshow(image_permuted[:patch_size, :, :]);
```




```
# Setup hyperparameters and make sure img_size and patch_size are
compatible
img_size = 224
patch_size = 16
num_patches = img_size/patch_size
assert img_size % patch_size == 0, "Image size must be divisible by
patch size"
print(f"Number of patches per row: {num_patches}\nPatch size:
{patch_size} pixels x {patch_size} pixels")

# Create a series of subplots
fig, axs = plt.subplots(nrows=1,
                        ncols=img_size // patch_size, # one column for
each patch

                        figsize=(num_patches, num_patches),
                        sharex=True,
                        sharey=True)

# Iterate through number of patches in the top row
for i, patch in enumerate(range(0, img_size, patch_size)):
    axs[i].imshow(image_permuted[:patch_size,
patch:patch+patch_size, :]); # keep height index constant, alter the
width index
    axs[i].set_xlabel(i+1) # set the label
    axs[i].set_xticks([])
    axs[i].set_yticks([])
```

Number of patches per row: 14.0
Patch size: 16 pixels x 16 pixels



```
# Setup hyperparameters and make sure img_size and patch_size are
compatible
img_size = 224
patch_size = 16
num_patches = img_size/patch_size
assert img_size % patch_size == 0, "Image size must be divisible by
patch size"
print(f"Number of patches per row: {num_patches}\
\nNumber of patches per column: {num_patches}\
\nTotal patches: {num_patches*num_patches}\
```

```

    \nPatch size: {patch_size} pixels x {patch_size} pixels")

# Create a series of subplots
fig, axs = plt.subplots(nrows=img_size // patch_size, # need int not
float
                        ncols=img_size // patch_size,
                        figsize=(num_patches, num_patches),
                        sharex=True,
                        sharey=True)

# Loop through height and width of image
for i, patch_height in enumerate(range(0, img_size, patch_size)): #
iterate through height
    for j, patch_width in enumerate(range(0, img_size, patch_size)): #
iterate through width

        # Plot the permuted image patch (image_permuted -> (Height,
Width, Color Channels))
        axs[i,
j].imshow(image_permuted[patch_height:patch_height+patch_size, #
iterate through height
patch_width:patch_width+patch_size, # iterate through width
:])) # get all color channels


        # Set up label information, remove the ticks for clarity and
set labels to outside
        axs[i, j].set_ylabel(i+1,
                                rotation="horizontal",
                                horizontalalignment="right",
                                verticalalignment="center")
        axs[i, j].set_xlabel(j+1)
        axs[i, j].set_xticks([])
        axs[i, j].set_yticks([])
        axs[i, j].label_outer()

# Set a super title
fig.suptitle(f"{class_names[label]} -> Patchified", fontsize=16)
plt.show()

Number of patches per row: 14.0
Number of patches per column: 14.0
Total patches: 196.0
Patch size: 16 pixels x 16 pixels

```


sushi -> Patchified



Membuat patch gambar dengan torch.nn.Conv2d()

```
from torch import nn
```

```
# Set the patch size  
patch_size=16
```

```
# Create the Conv2d layer with hyperparameters from the ViT paper  
conv2d = nn.Conv2d(in_channels=3, # number of color channels
```



```

        out_channels=768, # from Table 1: Hidden size D,
        this is the embedding size
        kernel_size=patch_size, # could also use
        (patch_size, patch_size)
        stride=patch_size,
        padding=0)

# View single image
plt.imshow(image.permute(1, 2, 0)) # adjust for matplotlib
plt.title(class_names[label])
plt.axis(False);

```

sushi



```

# Pass the image through the convolutional layer
image_out_of_conv = conv2d(image.unsqueeze(0)) # add a single batch
dimension (height, width, color_channels) -> (batch, height, width,
color_channels)
print(image_out_of_conv.shape)

torch.Size([1, 768, 14, 14])


# Plot random 5 convolutional feature maps
import random
random_indexes = random.sample(range(0, 758), k=5) # pick 5 numbers
between 0 and the embedding size
print(f"Showing random convolutional feature maps from indexes:
{random_indexes}")

# Create plot
fig, axs = plt.subplots(nrows=1, ncols=5, figsize=(12, 12))

```

```
# Plot random image feature maps
for i, idx in enumerate(random_indexes):
    image_conv_feature_map = image_out_of_conv[:, idx, :, :] # index
    on the output tensor of the convolutional layer
    axs[i].imshow(image_conv_feature_map.squeeze().detach().numpy())
    axs[i].set(xticklabels=[], yticklabels=[], xticks=[], yticks=[]);
```

Showing random convolutional feature maps from indexes: [571, 727, 734, 380, 90]



```
# Get a single feature map in tensor form
single_feature_map = image_out_of_conv[:, 0, :, :]
single_feature_map.requires_grad_()

(tensor([[[ 0.4732,  0.3567,  0.3377,  0.3736,  0.3208,  0.3913,
 0.3464,
           0.3702,  0.2541,  0.3594,  0.1984,  0.3982,  0.3741,
 0.1251],
          [ 0.4178,  0.4771,  0.3374,  0.3353,  0.3159,  0.4008,
 0.3448,
           0.3345,  0.5850,  0.4115,  0.2969,  0.2751,  0.6150,
 0.4188],
          [ 0.3209,  0.3776,  0.4970,  0.4272,  0.3301,  0.4787,
 0.2754,
           0.3726,  0.3298,  0.4631,  0.3087,  0.4915,  0.4129,
 0.4592],
          [ 0.4540,  0.4930,  0.5570,  0.2660,  0.2150,  0.2044,
 0.2766,
           0.2076,  0.3278,  0.3727,  0.2637,  0.2493,  0.2782,
 0.3664],
          [ 0.4920,  0.5671,  0.3298,  0.2992,  0.1437,  0.1701,
 0.1554,
           0.1375,  0.1377,  0.3141,  0.2694,  0.2771,  0.2412,
 0.3700],
          [ 0.5783,  0.5790,  0.4229,  0.5032,  0.1216,  0.1000,
 0.0356,
           0.1258, -0.0023,  0.1640,  0.2809,  0.2418,  0.2606,
 0.3787],
          [ 0.5334,  0.5645,  0.4781,  0.3307,  0.2391,  0.0461,
 0.0095,
           0.0542,  0.1012,  0.1331,  0.2446,  0.2526,  0.3323,
```

```

0.4120],
      [ 0.5724,  0.2840,  0.5188,  0.3934,  0.1328,  0.0776,
0.0235,
      0.1366,  0.3149,  0.2200,  0.2793,  0.2351,  0.4722,
0.4785],
      [ 0.4009,  0.4570,  0.4972,  0.5785,  0.2261,  0.1447, -
0.0028,
      0.2772,  0.2697,  0.4008,  0.3606,  0.3372,  0.4535,
0.4492],
      [ 0.5678,  0.5870,  0.5824,  0.3438,  0.5113,  0.0757,
0.1772,
      0.3677,  0.3572,  0.3742,  0.3820,  0.4868,  0.3781,
0.4694],
      [ 0.5845,  0.5877,  0.5826,  0.3212,  0.5276,  0.4840,
0.4825,
      0.5523,  0.5308,  0.5085,  0.5606,  0.5720,  0.4928,
0.5581],
      [ 0.5853,  0.5849,  0.5793,  0.3410,  0.4428,  0.4044,
0.3275,
      0.4958,  0.4366,  0.5750,  0.5494,  0.5868,  0.5557,
0.5069],
      [ 0.5880,  0.5888,  0.5796,  0.3377,  0.2635,  0.2347,
0.3145,
      0.3486,  0.5158,  0.5722,  0.5347,  0.5753,  0.5816,
0.4378],
      [ 0.5692,  0.5843,  0.5721,  0.5081,  0.2694,  0.2032,
0.1589,
      0.3464,  0.5349,  0.5768,  0.5739,  0.5764,  0.5394,
0.4482]]],
      grad_fn=<SliceBackward0>),
      True)

```

Meratakan penyematan tambalan dengan torch.nn.Flatten()

```

# Current tensor shape
print(f"Current tensor shape: {image_out_of_conv.shape} -> [batch,
embedding_dim, feature_map_height, feature_map_width]")

Current tensor shape: torch.Size([1, 768, 14, 14]) -> [batch,
embedding_dim, feature_map_height, feature_map_width]

# Create flatten layer
flatten = nn.Flatten(start_dim=2, # flatten feature_map_height
(dimension 2)
                        end_dim=3) # flatten feature_map_width (dimension
3)

# 1. View single image
plt.imshow(image.permute(1, 2, 0)) # adjust for matplotlib
plt.title(class_names[label])

```

```
plt.axis(False);
print(f"Original image shape: {image.shape}")

# 2. Turn image into feature maps
image_out_of_conv = conv2d(image.unsqueeze(0)) # add batch dimension
to avoid shape errors
print(f"Image feature map shape: {image_out_of_conv.shape}")

# 3. Flatten the feature maps
image_out_of_conv_flattened = flatten(image_out_of_conv)
print(f"Flattened image feature map shape:
{image_out_of_conv_flattened.shape}")

Original image shape: torch.Size([3, 224, 224])
Image feature map shape: torch.Size([1, 768, 14, 14])
Flattened image feature map shape: torch.Size([1, 768, 196])
```

sushi



```
# Get flattened image patch embeddings in right shape
image_out_of_conv_flattened_reshaped =
image_out_of_conv_flattened.permute(0, 2, 1) # [batch_size, P^2•C, N]
-> [batch_size, N, P^2•C]
print(f"Patch embedding sequence shape:
{image_out_of_conv_flattened_reshaped.shape} -> [batch_size,
num_patches, embedding_size]")

Patch embedding sequence shape: torch.Size([1, 196, 768]) ->
[batch_size, num_patches, embedding_size]

# Get a single flattened feature map
single_flattened_feature_map = image_out_of_conv_flattened_reshaped[:,
```

```

[:, 0] # index: (batch_size, number_of_patches, embedding_dimension)

# Plot the flattened feature map visually
plt.figure(figsize=(22, 22))
plt.imshow(single_flattened_feature_map.detach().numpy())
plt.title(f"Flattened feature map shape:
{single_flattened_feature_map.shape}")
plt.axis(False);

```

Flattened feature map shape: torch.Size([1, 196])

```

# See the flattened feature map as a tensor
single_flattened_feature_map,
single_flattened_feature_map.requires_grad,
single_flattened_feature_map.shape

(tensor([[ 0.4732,  0.3567,  0.3377,  0.3736,  0.3208,  0.3913,
 0.3464,  0.3702,
          0.2541,  0.3594,  0.1984,  0.3982,  0.3741,  0.1251,
 0.4178,  0.4771,
          0.3374,  0.3353,  0.3159,  0.4008,  0.3448,  0.3345,
 0.5850,  0.4115,
          0.2969,  0.2751,  0.6150,  0.4188,  0.3209,  0.3776,
 0.4970,  0.4272,
          0.3301,  0.4787,  0.2754,  0.3726,  0.3298,  0.4631,
 0.3087,  0.4915,
          0.4129,  0.4592,  0.4540,  0.4930,  0.5570,  0.2660,
 0.2150,  0.2044,
          0.2766,  0.2076,  0.3278,  0.3727,  0.2637,  0.2493,
 0.2782,  0.3664,
          0.4920,  0.5671,  0.3298,  0.2992,  0.1437,  0.1701,
 0.1554,  0.1375,
          0.1377,  0.3141,  0.2694,  0.2771,  0.2412,  0.3700,
 0.5783,  0.5790,
          0.4229,  0.5032,  0.1216,  0.1000,  0.0356,  0.1258, -
 0.0023,  0.1640,
          0.2809,  0.2418,  0.2606,  0.3787,  0.5334,  0.5645,
 0.4781,  0.3307,
          0.2391,  0.0461,  0.0095,  0.0542,  0.1012,  0.1331,
 0.2446,  0.2526,
          0.3323,  0.4120,  0.5724,  0.2840,  0.5188,  0.3934,
 0.1328,  0.0776,
          0.0235,  0.1366,  0.3149,  0.2200,  0.2793,  0.2351,
 0.4722,  0.4785,
          0.4009,  0.4570,  0.4972,  0.5785,  0.2261,  0.1447, -
 0.0028,  0.2772,
          0.2697,  0.4008,  0.3606,  0.3372,  0.4535,  0.4492,
 0.5678,  0.5870,
          0.5824,  0.3438,  0.5113,  0.0757,  0.1772,  0.3677,

```

```

0.3572, 0.3742,
        0.3820, 0.4868, 0.3781, 0.4694, 0.5845, 0.5877,
0.5826, 0.3212,
        0.5276, 0.4840, 0.4825, 0.5523, 0.5308, 0.5085,
0.5606, 0.5720,
        0.4928, 0.5581, 0.5853, 0.5849, 0.5793, 0.3410,
0.4428, 0.4044,
        0.3275, 0.4958, 0.4366, 0.5750, 0.5494, 0.5868,
0.5557, 0.5069,
        0.5880, 0.5888, 0.5796, 0.3377, 0.2635, 0.2347,
0.3145, 0.3486,
        0.5158, 0.5722, 0.5347, 0.5753, 0.5816, 0.4378,
0.5692, 0.5843,
        0.5721, 0.5081, 0.2694, 0.2032, 0.1589, 0.3464,
0.5349, 0.5768,
        0.5739, 0.5764, 0.5394, 0.4482]],
grad_fn=<SelectBackward0>),
    True,
    torch.Size([1, 196]))

```

Mengubah lapisan penyematanan patch ViT menjadi modul PyTorch

```

# 1. Create a class which subclasses nn.Module
class PatchEmbedding(nn.Module):
    """Turns a 2D input image into a 1D sequence learnable embedding
    vector.

    Args:
        in_channels (int): Number of color channels for the input
        images. Defaults to 3.
        patch_size (int): Size of patches to convert input image into.
        Defaults to 16.
        embedding_dim (int): Size of embedding to turn image into.
        Defaults to 768.
    """
    # 2. Initialize the class with appropriate variables
    def __init__(self,
                 in_channels:int=3,
                 patch_size:int=16,
                 embedding_dim:int=768):
        super().__init__()

    # 3. Create a layer to turn an image into patches
    self.patcher = nn.Conv2d(in_channels=in_channels,
                             out_channels=embedding_dim,
                             kernel_size=patch_size,
                             stride=patch_size,
                             padding=0)

    # 4. Create a layer to flatten the patch feature maps into a

```

```

single dimension
    self.flatten = nn.Flatten(start_dim=2, # only flatten the
feature map dimensions into a single vector
                                end_dim=3)

# 5. Define the forward method
def forward(self, x):
    # Create assertion to check that inputs are the correct shape
    image_resolution = x.shape[-1]
    assert image_resolution % patch_size == 0, f"Input image size
must be divisble by patch size, image shape: {image_resolution}, patch
size: {patch_size}"

    # Perform the forward pass
    x_patched = self.patcher(x)
    x_flattened = self.flatten(x_patched)
    # 6. Make sure the output shape has the right order
    return x_flattened.permute(0, 2, 1) # adjust so the embedding
is on the final dimension [batch_size, P^2*C, N] -> [batch_size, N,
P^2*C]

set_seeds()

# Create an instance of patch embedding layer
patchify = PatchEmbedding(in_channels=3,
                           patch_size=16,
                           embedding_dim=768)

# Pass a single image through
print(f"Input image shape: {image.unsqueeze(0).shape}")
patch_embedded_image = patchify(image.unsqueeze(0)) # add an extra
batch dimension on the 0th index, otherwise will error
print(f"Output patch embedding shape: {patch_embedded_image.shape}")

Input image shape: torch.Size([1, 3, 224, 224])
Output patch embedding shape: torch.Size([1, 196, 768])

# Create random input sizes
random_input_image = (1, 3, 224, 224)
random_input_image_error = (1, 3, 250, 250) # will error because image
size is incompatible with patch_size

# # Get a summary of the input and outputs of PatchEmbedding
(uncomment for full output)
# summary(PatchEmbedding(),
#         input_size=random_input_image, # try swapping this for
"random_input_image_error"
#         col_names=["input_size", "output_size", "num_params",
"trainable"],

```



```
# col_width=20,  
# row_settings=["var_names"])
```

Membuat penyematan token kelas

```
# View the patch embedding and patch embedding shape
print(patch_embedded_image)
print(f"Patch embedding shape: {patch_embedded_image.shape} -> [batch_size, number_of_patches, embedding_dimension]")

tensor([[[[-0.9145,  0.2454, -0.2292, ...,  0.6768, -0.4515,  0.3496],
          [-0.7427,  0.1955, -0.3570, ...,  0.5823, -0.3458,  0.3261],
          [-0.7589,  0.2633, -0.1695, ...,  0.5897, -0.3980,  0.0761],
          ..., 
          [-1.0072,  0.2795, -0.2804, ...,  0.7624, -0.4584,  0.3581],
          [-0.9839,  0.1652, -0.1576, ...,  0.7489, -0.5478,  0.3486],
          [-0.9260,  0.1383, -0.1157, ...,  0.5847, -0.4717,  0.3112]]]],
        grad_fn=<PermuteBackward0>)
Patch embedding shape: torch.Size([1, 196, 768]) -> [batch_size, number_of_patches, embedding_dimension]

# Get the batch size and embedding dimension
batch_size = patch_embedded_image.shape[0]
embedding_dimension = patch_embedded_image.shape[-1]

# Create the class token embedding as a learnable parameter that shares the same size as the embedding dimension (D)
class_token = nn.Parameter(torch.ones(batch_size, 1, embedding_dimension), # [batch_size, number_of_tokens, embedding_dimension]
                             requires_grad=True) # make sure the embedding is learnable

# Show the first 10 examples of the class_token
print(class_token[:, :, :10])

# Print the class_token shape
print(f"Class token shape: {class_token.shape} -> [batch_size, number_of_tokens, embedding_dimension]")

tensor([[[[1., 1., 1., 1., 1., 1., 1., 1., 1., 1.] ]]],
        grad_fn=<SliceBackward0>)
Class token shape: torch.Size([1, 1, 768]) -> [batch_size, number_of_tokens, embedding_dimension]

# Add the class token embedding to the front of the patch embedding
patch_embedded_image_with_class_embedding = torch.cat((class_token, patch_embedded_image), dim=1) # concat
```

on first dimension

Print the sequence of patch embeddings with the prepended class token embedding

```
print(patch_embedded_image_with_class_embedding)
print(f"Sequence of patch embeddings with class token prepended shape:
{patch_embedded_image_with_class_embedding.shape} -> [batch_size,
number_of_patches, embedding_dimension]")
```

```
tensor([[[ 1.0000,  1.0000,  1.0000, ...,  1.0000,  1.0000,  1.0000],
          [-0.9145,  0.2454, -0.2292, ...,  0.6768, -0.4515,  0.3496],
          [-0.7427,  0.1955, -0.3570, ...,  0.5823, -0.3458,  0.3261],
          ...,
          [-1.0072,  0.2795, -0.2804, ...,  0.7624, -0.4584,  0.3581],
          [-0.9839,  0.1652, -0.1576, ...,  0.7489, -0.5478,  0.3486],
          [-0.9260,  0.1383, -0.1157, ...,  0.5847, -0.4717,
0.3112]]],
```

```
      grad_fn=<CatBackward0>)
```

```
Sequence of patch embeddings with class token prepended shape:
torch.Size([1, 197, 768]) -> [batch_size, number_of_patches,
embedding_dimension]
```

Membuat posisi penyematan

View the sequence of patch embeddings with the prepended class embedding

```
patch_embedded_image_with_class_embedding,
patch_embedded_image_with_class_embedding.shape
```

```
(tensor([[[ 1.0000,  1.0000,  1.0000, ...,  1.0000,  1.0000,
1.0000],
          [-0.9145,  0.2454, -0.2292, ...,  0.6768, -0.4515,
0.3496],
          [-0.7427,  0.1955, -0.3570, ...,  0.5823, -0.3458,
0.3261],
          ...,
          [-1.0072,  0.2795, -0.2804, ...,  0.7624, -0.4584,
0.3581],
          [-0.9839,  0.1652, -0.1576, ...,  0.7489, -0.5478,
0.3486],
          [-0.9260,  0.1383, -0.1157, ...,  0.5847, -0.4717,
0.3112]]],
```

```
      grad_fn=<CatBackward0>),
```

```
torch.Size([1, 197, 768]))
```

Calculate N (number of patches)

```
number_of_patches = int((height * width) / patch_size**2)
```

Get embedding dimension

```
embedding_dimension =
```

```

patch_embedded_image_with_class_embedding.shape[2]

# Create the learnable 1D position embedding
position_embedding = nn.Parameter(torch.ones(1,
                                              number_of_patches+1,
                                              embedding_dimension),
                                  requires_grad=True) # make sure it's
learnable

# Show the first 10 sequences and 10 position embedding values and
check the shape of the position embedding
print(position_embedding[:, :10, :10])
print(f"Position embeddding shape: {position_embedding.shape} ->
[batch_size, number_of_patches, embedding_dimension]")

tensor([[[[1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.],
          [1., 1., 1., 1., 1., 1., 1., 1., 1., 1.]]],
        grad_fn=<SliceBackward0>)]
Position embeddding shape: torch.Size([1, 197, 768]) -> [batch_size,
number_of_patches, embedding_dimension]

# Add the position embedding to the patch and class token embedding
patch_and_position_embedding =
patch_embedded_image_with_class_embedding + position_embedding
print(patch_and_position_embedding)
print(f"Patch embeddings, class token prepended and positional
embeddings added shape: {patch_and_position_embedding.shape} ->
[batch_size, number_of_patches, embedding_dimension]")

tensor([[[[ 2.0000,  2.0000,  2.0000, ...,  2.0000,  2.0000,  2.0000],
          [ 0.0855,  1.2454,  0.7708, ...,  1.6768,  0.5485,  1.3496],
          [ 0.2573,  1.1955,  0.6430, ...,  1.5823,  0.6542,  1.3261],
          ...,
          [-0.0072,  1.2795,  0.7196, ...,  1.7624,  0.5416,  1.3581],
          [ 0.0161,  1.1652,  0.8424, ...,  1.7489,  0.4522,  1.3486],
          [ 0.0740,  1.1383,  0.8843, ...,  1.5847,  0.5283,
1.3112]]],
        grad_fn=<AddBackward0>)]
Patch embeddings, class token prepended and positional embeddings
added shape: torch.Size([1, 197, 768]) -> [batch_size,
number_of_patches, embedding_dimension]

```

Menyatukan semuanya: dari gambar hingga penyematan

```
set_seeds()

# 1. Set patch size
patch_size = 16

# 2. Print shape of original image tensor and get the image dimensions
print(f"Image tensor shape: {image.shape}")
height, width = image.shape[1], image.shape[2]

# 3. Get image tensor and add batch dimension
x = image.unsqueeze(0)
print(f"Input image with batch dimension shape: {x.shape}")

# 4. Create patch embedding layer
patch_embedding_layer = PatchEmbedding(in_channels=3,
                                       patch_size=patch_size,
                                       embedding_dim=768)

# 5. Pass image through patch embedding layer
patch_embedding = patch_embedding_layer(x)
print(f"Patching embedding shape: {patch_embedding.shape}")

# 6. Create class token embedding
batch_size = patch_embedding.shape[0]
embedding_dimension = patch_embedding.shape[-1]
class_token = nn.Parameter(torch.ones(batch_size, 1,
                                       embedding_dimension),
                           requires_grad=True) # make sure it's
learnable
print(f"Class token embedding shape: {class_token.shape}")

# 7. Prepend class token embedding to patch embedding
patch_embedding_class_token = torch.cat((class_token,
                                         patch_embedding), dim=1)
print(f"Patch embedding with class token shape:
{patch_embedding_class_token.shape}")

# 8. Create position embedding
number_of_patches = int((height * width) / patch_size**2)
position_embedding = nn.Parameter(torch.ones(1, number_of_patches+1,
                                             embedding_dimension),
                                   requires_grad=True) # make sure it's
learnable

# 9. Add position embedding to patch embedding with class token
patch_and_position_embedding = patch_embedding_class_token +
position_embedding
print(f"Patch and position embedding shape:
{patch_and_position_embedding.shape}")
```

```

Image tensor shape: torch.Size([3, 224, 224])
Input image with batch dimension shape: torch.Size([1, 3, 224, 224])
Patching embedding shape: torch.Size([1, 196, 768])
Class token embedding shape: torch.Size([1, 1, 768])
Patch embedding with class token shape: torch.Size([1, 197, 768])
Patch and position embedding shape: torch.Size([1, 197, 768])

# 1. Create a class that inherits from nn.Module
class MultiheadSelfAttentionBlock(nn.Module):
    """Creates a multi-head self-attention block ("MSA block" for
    short).
    """
    # 2. Initialize the class with hyperparameters from Table 1
    def __init__(self,
        embedding_dim:int=768, # Hidden size D from Table 1
        num_heads:int=12, # Heads from Table 1 for ViT-Base
        attn_dropout:float=0): # doesn't look like the paper
        # uses any dropout in MSABlocks
        super().__init__()

        # 3. Create the Norm layer (LN)
        self.layer_norm = nn.LayerNorm(normalized_shape=embedding_dim)

        # 4. Create the Multi-Head Attention (MSA) layer
        self.multihead_attn =
nn.MultiheadAttention(embed_dim=embedding_dim,
num_heads=num_heads,
dropout=attn_dropout,
batch_first=True)
# does our batch dimension come first?

    # 5. Create a forward() method to pass the data througuh the layers
    def forward(self, x):
        x = self.layer_norm(x)
        attn_output, _ = self.multihead_attn(query=x, # query
embeddings
key=x, # key embeddings
value=x, # value
embeddings
need_weights=False) # do
we need the weights or just the layer outputs?
        return attn_output

# Create an instance of MSABlock
multihead_self_attention_block =
MultiheadSelfAttentionBlock(embedding_dim=768, # from Table 1

```

```

num_heads=12) # from Table 1

# Pass patch and position image embedding through MSABlock
patched_image_through_msa_block =
multihead_self_attention_block(patch_and_position_embedding)
print(f"Input shape of MSA block:
{patch_and_position_embedding.shape}")
print(f"Output shape MSA block:
{patched_image_through_msa_block.shape}")

Input shape of MSA block: torch.Size([1, 197, 768])
Output shape MSA block: torch.Size([1, 197, 768])

# 1. Create a class that inherits from nn.Module
class MLPBlock(nn.Module):
    """Creates a layer normalized multilayer perceptron block ("MLP
    block" for short)."""
    # 2. Initialize the class with hyperparameters from Table 1 and
    Table 3
    def __init__(self,
        embedding_dim:int=768, # Hidden Size D from Table 1
        mlp_size:int=3072, # MLP size from Table 1 for ViT-
        dropout:float=0.1): # Dropout from Table 3 for ViT-

        super().__init__()

        # 3. Create the Norm layer (LN)
        self.layer_norm = nn.LayerNorm(normalized_shape=embedding_dim)

        # 4. Create the Multilayer perceptron (MLP) layer(s)
        self.mlp = nn.Sequential(
            nn.Linear(in_features=embedding_dim,
                out_features=mlp_size),
            nn.GELU(), # "The MLP contains two layers with a GELU non-
            linearity (section 3.1)."
            nn.Dropout(p=dropout),
            nn.Linear(in_features=mlp_size, # needs to take same
            in_features as out_features of layer above
                out_features=embedding_dim), # take back to
            embedding_dim
            nn.Dropout(p=dropout) # "Dropout, when used, is applied
            after every dense layer.."
        )

        # 5. Create a forward() method to pass the data throug the layers
        def forward(self, x):
            x = self.layer_norm(x)

```

```

        x = self.mlp(x)
        return x

# Create an instance of MLPBlock
mlp_block = MLPBlock(embedding_dim=768, # from Table 1
                      mlp_size=3072, # from Table 1
                      dropout=0.1) # from Table 3

# Pass output of MSABlock through MLPBlock
patched_image_through_mlp_block =
mlp_block(patch_image_through_msa_block)
print(f"Input shape of MLP block:
{patched_image_through_msa_block.shape}")
print(f"Output shape MLP block:
{patched_image_through_mlp_block.shape}")

Input shape of MLP block: torch.Size([1, 197, 768])
Output shape MLP block: torch.Size([1, 197, 768])

```

Buat Encoder Transformatior

Membuat Transformer Encoder dengan menggabungkan lapisan yang kami buat khusus

```

# 1. Create a class that inherits from nn.Module
class TransformerEncoderBlock(nn.Module):
    """Creates a Transformer Encoder block."""
    # 2. Initialize the class with hyperparameters from Table 1 and
    Table 3
    def __init__(self,
                  embedding_dim:int=768, # Hidden size D from Table 1
                  num_heads:int=12, # Heads from Table 1 for ViT-Base
                  mlp_size:int=3072, # MLP size from Table 1 for ViT-
                  Base
                  mlp_dropout:float=0.1, # Amount of dropout for dense
                  layers from Table 3 for ViT-Base
                  attn_dropout:float=0): # Amount of dropout for
                  attention layers
        super().__init__()

        # 3. Create MSA block (equation 2)
        self.msa_block =
MultiheadSelfAttentionBlock(embedding_dim=embedding_dim,
num_heads=num_heads,
attn_dropout=attn_dropout)

```



```

# 4. Create MLP block (equation 3)
self.mlp_block = MLPBlock(embedding_dim=embedding_dim,
                           mlp_size=mlp_size,
                           dropout=mlp_dropout)

# 5. Create a forward() method
def forward(self, x):

    # 6. Create residual connection for MSA block (add the input
    to the output)
    x = self.msa_block(x) + x

    # 7. Create residual connection for MLP block (add the input
    to the output)
    x = self.mlp_block(x) + x

    return x

# Create an instance of TransformerEncoderBlock
transformer_encoder_block = TransformerEncoderBlock()

# # Print an input and output summary of our Transformer Encoder
# (uncomment for full output)
# summary(model=transformer_encoder_block,
#         input_size=(1, 197, 768), # (batch_size, num_patches,
# embedding_dimension)
#         col_names=["input_size", "output_size", "num_params",
# "trainable"],
#         col_width=20,
#         row_settings=["var_names"])

```

Membuat Transformer Encoder dengan lapisan Transformer PyTorch

```

# Create the same as above with torch.nn.TransformerEncoderLayer()
torch_transformer_encoder_layer =
nn.TransformerEncoderLayer(d_model=768, # Hidden size D from Table 1
for ViT-Base
nhead=12,

# Heads from Table 1 for ViT-Base

dim_feedforward=3072, # MLP size from Table 1 for ViT-Base

dropout=0.1, # Amount of dropout for dense layers from Table 3 for
ViT-Base

activation="gelu", # GELU non-linear activation

batch_first=True, # Do our batches come first?

norm_first=True) # Normalize first or after MSA/MLP layers?

```

```


torch_transformer_encoder_layer
TransformerEncoderLayer(
  (self_attn): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
  )
  (linear1): Linear(in_features=768, out_features=3072, bias=True)
  (dropout): Dropout(p=0.1, inplace=False)
  (linear2): Linear(in_features=3072, out_features=768, bias=True)
  (norm1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
  (norm2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
  (dropout1): Dropout(p=0.1, inplace=False)
  (dropout2): Dropout(p=0.1, inplace=False)
)

# # Get the output of PyTorch's version of the Transformer Encoder
(uncomment for full output)
# summary(model=torch_transformer_encoder_layer,
#         input_size=(1, 197, 768), # (batch_size, num_patches,
embedding_dimension)
#         col_names=["input_size", "output_size", "num_params",
"trainable"],
#         col_width=20,
#         row_settings=["var_names"])

```

Menggabungkan semuanya untuk menciptakan ViT

```

# 1. Create a ViT class that inherits from nn.Module
class ViT(nn.Module):
    """Creates a Vision Transformer architecture with ViT-Base
hyperparameters by default."""
    # 2. Initialize the class with hyperparameters from Table 1 and
Table 3
    def __init__(self,

        img_size:int=224, # Training resolution from Table 3
        in_channels:int=3, # Number of channels in input
        patch_size:int=16, # Patch size
        num_transformer_layers:int=12, # Layers from Table 1
        embedding_dim:int=768, # Hidden size D from Table 1
        mlp_size:int=3072, # MLP size from Table 1 for ViT-
        num_heads:int=12, # Heads from Table 1 for ViT-Base
        attn_dropout:float=0, # Dropout for attention
        projection
    )

```

```

        mlp_dropout:float=0.1, # Dropout for dense/MLP layers
        embedding_dropout:float=0.1, # Dropout for patch and
position embeddings
        num_classes:int=1000): # Default for ImageNet but can
customize this
    super().__init__() # don't forget the super().__init__()!

    # 3. Make the image size is divisble by the patch size
    assert img_size % patch_size == 0, f"Image size must be
divisible by patch size, image size: {img_size}, patch size:
{patch_size}."

    # 4. Calculate number of patches (height * width/patch^2)
    self.num_patches = (img_size * img_size) // patch_size**2

    # 5. Create learnable class embedding (needs to go at front of
sequence of patch embeddings)
    self.class_embedding = nn.Parameter(data=torch.randn(1, 1,
embedding_dim),
                                       requires_grad=True)

    # 6. Create learnable position embedding
    self.position_embedding = nn.Parameter(data=torch.randn(1,
self.num_patches+1, embedding_dim),
                                       requires_grad=True)

    # 7. Create embedding dropout value
    self.embedding_dropout = nn.Dropout(p=embedding_dropout)

    # 8. Create patch embedding layer
    self.patch_embedding = PatchEmbedding(in_channels=in_channels,
                                       patch_size=patch_size,
embedding_dim=embedding_dim)

    # 9. Create Transformer Encoder blocks (we can stack
Transformer Encoder blocks using nn.Sequential())
    # Note: The "*" means "all"
    self.transformer_encoder =
nn.Sequential(*[TransformerEncoderBlock(embedding_dim=embedding_dim,
num_heads=num_heads,
mlp_size=mlp_size,
mlp_dropout=mlp_dropout) for _ in range(num_transformer_layers)])

    # 10. Create classifier head
    self.classifier = nn.Sequential(
        nn.LayerNorm(normalized_shape=embedding_dim),

```

```

        nn.Linear(in_features=embedding_dim,
                  out_features=num_classes)
    )

    # 11. Create a forward() method
    def forward(self, x):

        # 12. Get batch size
        batch_size = x.shape[0]

        # 13. Create class token embedding and expand it to match the
        batch size (equation 1)
        class_token = self.class_embedding.expand(batch_size, -1, -1)
        # "-1" means to infer the dimension (try this line on its own)

        # 14. Create patch embedding (equation 1)
        x = self.patch_embedding(x)

        # 15. Concat class embedding and patch embedding (equation 1)
        x = torch.cat((class_token, x), dim=1)

        # 16. Add position embedding to patch embedding (equation 1)
        x = self.position_embedding + x

        # 17. Run embedding dropout (Appendix B.1)
        x = self.embedding_dropout(x)

        # 18. Pass patch, position and class embedding through
        transformer encoder layers (equations 2 & 3)
        x = self.transformer_encoder(x)

        # 19. Put 0 index logit through classifier (equation 4)
        x = self.classifier(x[:, 0]) # run on each sample in a batch
        at 0 index

        return x

# Example of creating the class embedding and expanding over a batch
dimension
batch_size = 32
class_token_embedding_single = nn.Parameter(data=torch.randn(1, 1,
768)) # create a single learnable class token
class_token_embedding_expanded =
class_token_embedding_single.expand(batch_size, -1, -1) # expand the
single learnable class token across the batch dimension, "-1" means to
"infer the dimension"

# Print out the change in shapes
print(f"Shape of class token embedding single:
{class_token_embedding_single.shape}")

```

```

print(f"Shape of class token embedding expanded:
{class_token_embedding_expanded.shape}")

Shape of class token embedding single: torch.Size([1, 1, 768])
Shape of class token embedding expanded: torch.Size([32, 1, 768])

set_seeds()

# Create a random tensor with same shape as a single image
random_image_tensor = torch.randn(1, 3, 224, 224) # (batch_size,
color_channels, height, width)

# Create an instance of ViT with the number of classes we're working
with (pizza, steak, sushi)
vit = ViT(num_classes=len(class_names))

# Pass the random image tensor to our ViT instance
vit(random_image_tensor)

tensor([[ -0.2377,  0.7360,  1.2137]], grad_fn=<AddmmBackward0>)

```

Mendapatkan ringkasan visual model ViT kami

```

from torchinfo import summary

# # Print a summary of our custom ViT model using torchinfo (uncomment
for actual output)
# summary(model=vit,
#         input_size=(32, 3, 224, 224), # (batch_size, color_channels,
height, width)
#         # col_names=["input_size"], # uncomment for smaller output
#         col_names=["input_size", "output_size", "num_params",
"trainable"],
#         col_width=20,
#         row_settings=["var_names"]
# )

```

Membuat pengoptimal

Melatih model ViT kita

```

from going_modular.going_modular import engine

# Setup the optimizer to optimize our ViT model parameters using
hyperparameters from the ViT paper
optimizer = torch.optim.Adam(params=vit.parameters(),
                             lr=3e-3, # Base LR from Table 3 for ViT-*
ImageNet-1k
                             betas=(0.9, 0.999), # default values but
also mentioned in ViT paper section 4.1 (Training & Fine-tuning)

```

```

weight_decay=0.3) # from the ViT paper
section 4.1 (Training & Fine-tuning) and Table 3 for ViT-* ImageNet-1k

# Setup the loss function for multi-class classification
loss_fn = torch.nn.CrossEntropyLoss()

# Set the seeds
set_seeds()

# Train the model and save the training results to a dictionary
results = engine.train(model=vit,
                        train_dataloader=train_dataloader,
                        test_dataloader=test_dataloader,
                        optimizer=optimizer,
                        loss_fn=loss_fn,
                        epochs=10,
                        device=device)

{"model_id": "97484323a38248e98ded3df3e074655c", "version_major": 2, "version_minor": 0}

Epoch: 1 | train_loss: 4.8759 | train_acc: 0.2891 | test_loss: 1.0465
| test_acc: 0.5417
Epoch: 2 | train_loss: 1.5900 | train_acc: 0.2617 | test_loss: 1.5876
| test_acc: 0.1979
Epoch: 3 | train_loss: 1.4644 | train_acc: 0.2617 | test_loss: 1.2738
| test_acc: 0.1979
Epoch: 4 | train_loss: 1.3159 | train_acc: 0.2773 | test_loss: 1.7498
| test_acc: 0.1979
Epoch: 5 | train_loss: 1.3114 | train_acc: 0.3008 | test_loss: 1.7444
| test_acc: 0.2604
Epoch: 6 | train_loss: 1.2445 | train_acc: 0.3008 | test_loss: 1.9704
| test_acc: 0.1979
Epoch: 7 | train_loss: 1.2050 | train_acc: 0.3984 | test_loss: 3.5480
| test_acc: 0.1979
Epoch: 8 | train_loss: 1.4368 | train_acc: 0.4258 | test_loss: 1.8324
| test_acc: 0.2604
Epoch: 9 | train_loss: 1.5757 | train_acc: 0.2344 | test_loss: 1.2848
| test_acc: 0.5417
Epoch: 10 | train_loss: 1.4658 | train_acc: 0.4023 | test_loss: 1.2389
| test_acc: 0.2604

```

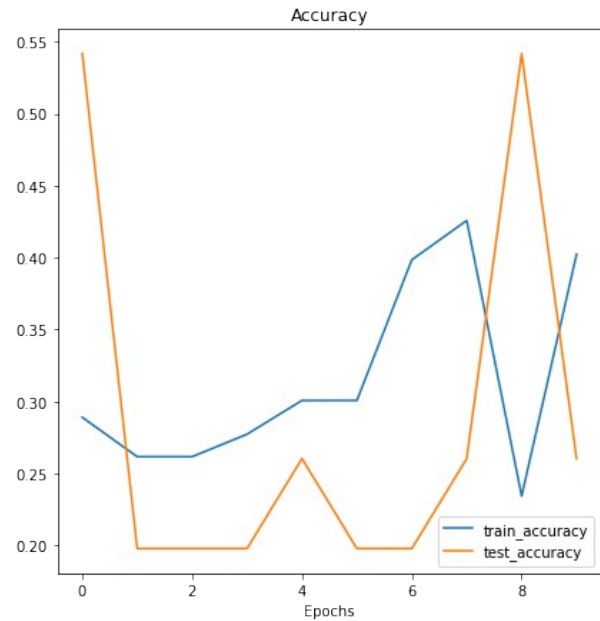

Plot kurva kerugian model ViT kita

```

from helper_functions import plot_loss_curves

# Plot our ViT model's loss curves
plot_loss_curves(results)

```



Mendapatkan model ViT terlatih dan membuat ekstraktor fitur

```
# The following requires torch v0.12+ and torchvision v0.13+
import torch
import torchvision
print(torch.__version__)
print(torchvision.__version__)

1.12.0+cu102
0.13.0+cu102
```

Then we'll setup device-agnostic code.

```
device = "cuda" if torch.cuda.is_available() else "cpu"
device

'cuda'

# 1. Get pretrained weights for ViT-Base
pretrained_vit_weights = torchvision.models.ViT_B_16_Weights.DEFAULT #
requires torchvision >= 0.13, "DEFAULT" means best available

# 2. Setup a ViT model instance with pretrained weights
pretrained_vit =
torchvision.models.vit_b_16(weights=pretrained_vit_weights).to(device)

# 3. Freeze the base parameters
for parameter in pretrained_vit.parameters():
    parameter.requires_grad = False

# 4. Change the classifier head (set the seeds to ensure same
```



```

initialization with linear head)
set_seeds()
pretrained_vit.heads = nn.Linear(in_features=768,
out_features=len(class_names)).to(device)
# pretrained_vit # uncomment for model output

# # Print a summary using torchinfo (uncomment for actual output)
# summary(model=pretrained_vit,
#         input_size=(32, 3, 224, 224), # (batch_size, color_channels,
height, width)
#         # col_names=["input_size"], # uncomment for smaller output
#         col_names=["input_size", "output_size", "num_params",
"trainable"],
#         col_width=20,
#         row_settings=["var_names"]
# )

```

Mempersiapkan data untuk model ViT yang telah dilatih sebelumnya

```

from helper_functions import download_data

# Download pizza, steak, sushi images from GitHub
image_path =
download_data(source="https://github.com/mrdbourke/pytorch-deep-
learning/raw/main/data/pizza_steak_sushi.zip",
              destination="pizza_steak_sushi")

image_path

[INFO] data/pizza_steak_sushi directory exists, skipping download.
PosixPath('data/pizza_steak_sushi')

# Setup train and test directory paths
train_dir = image_path / "train"
test_dir = image_path / "test"
train_dir, test_dir

(PosixPath('data/pizza_steak_sushi/train'),
 PosixPath('data/pizza_steak_sushi/test'))

# Get automatic transforms from pretrained ViT weights
pretrained_vit_transforms = pretrained_vit_weights.transforms()
print(pretrained_vit_transforms)

ImageClassification(
    crop_size=[224]
    resize_size=[256]
    mean=[0.485, 0.456, 0.406]
    std=[0.229, 0.224, 0.225]
    interpolation=InterpolationMode.BILINEAR
)

```

```
# Setup dataloaders
train_dataloader_pretrained, test_dataloader_pretrained, class_names =
data_setup.create_dataloaders(train_dir=train_dir,

test_dir=test_dir,

transform=pretrained_vit_transforms,

batch_size=32) # Could increase if we had more samples, such as here:
https://arxiv.org/abs/2205.01580 (there are other improvements there
too...)
```

Melatih model ViT ekstraktor fitur

```
from going_modular.going_modular import engine

# Create optimizer and loss function
optimizer = torch.optim.Adam(params=pretrained_vit.parameters(),
                              lr=1e-3)
loss_fn = torch.nn.CrossEntropyLoss()

# Train the classifier head of the pretrained ViT feature extractor
model
set_seeds()
pretrained_vit_results = engine.train(model=pretrained_vit,

train_dataloader=train_dataloader_pretrained,

test_dataloader=test_dataloader_pretrained,
                                optimizer=optimizer,
                                loss_fn=loss_fn,
                                epochs=10,
                                device=device)

{"model_id": "e47702187773418aafc32e0078ff1895", "version_major": 2, "version_minor": 0}


Epoch: 1 | train_loss: 0.7665 | train_acc: 0.7227 | test_loss: 0.5432
| test_acc: 0.8665
Epoch: 2 | train_loss: 0.3428 | train_acc: 0.9453 | test_loss: 0.3263
| test_acc: 0.8977
Epoch: 3 | train_loss: 0.2064 | train_acc: 0.9531 | test_loss: 0.2707
| test_acc: 0.9081
Epoch: 4 | train_loss: 0.1556 | train_acc: 0.9570 | test_loss: 0.2422
| test_acc: 0.9081
Epoch: 5 | train_loss: 0.1246 | train_acc: 0.9727 | test_loss: 0.2279
| test_acc: 0.8977
Epoch: 6 | train_loss: 0.1216 | train_acc: 0.9766 | test_loss: 0.2129
| test_acc: 0.9280
Epoch: 7 | train_loss: 0.0938 | train_acc: 0.9766 | test_loss: 0.2352
```

```
| test_acc: 0.8883
Epoch: 8 | train_loss: 0.0797 | train_acc: 0.9844 | test_loss: 0.2281
| test_acc: 0.8778
Epoch: 9 | train_loss: 0.1098 | train_acc: 0.9883 | test_loss: 0.2074
| test_acc: 0.9384
Epoch: 10 | train_loss: 0.0650 | train_acc: 0.9883 | test_loss: 0.1804
| test_acc: 0.9176
```

Plot kurva kerugian model ViT ekstraktor fitur

```
# Plot the loss curves
from helper_functions import plot_loss_curves

plot_loss_curves(pretrained_vit_results)
```



Simpan model ViT ekstraktor fitur dan periksa ukuran file

```
# Save the model
from going_modular.going_modular import utils

utils.save_model(model=pretrained_vit,
                  target_dir="models",

model_name="08_pretrained_vit_feature_extractor_pizza_steak_sushi.pth"
)

[INFO] Saving model to:
models/08_pretrained_vit_feature_extractor_pizza_steak_sushi.pth

from pathlib import Path
```

```

# Get the model size in bytes then convert to megabytes
pretrained_vit_model_size =
Path("models/08_pretrained_vit_feature_extractor_pizza_steak_sushi.pth
").stat().st_size // (1024*1024) # division converts bytes to
megabytes (roughly)
print(f"Pretrained ViT feature extractor model size:
{pretrained_vit_model_size} MB")

```

Pretrained ViT feature extractor model size: 327 MB

Buat prediksi pada gambar khusus

```

import requests

# Import function to make predictions on images and plot them
from going_modular.going_modular.predictions import
pred_and_plot_image

# Setup custom image path
custom_image_path = image_path / "04-pizza-dad.jpeg"

# Download the image if it doesn't already exist
if not custom_image_path.is_file():
    with open(custom_image_path, "wb") as f:
        # When downloading from GitHub, need to use the "raw" file
        link
        request =
requests.get("https://raw.githubusercontent.com/mrdbourke/pytorch-
deep-learning/main/images/04-pizza-dad.jpeg")
        print(f"Downloading {custom_image_path}...")
        f.write(request.content)
else:
    print(f"{custom_image_path} already exists, skipping download.")

# Predict on custom image
pred_and_plot_image(model=pretrained_vit,
                    image_path=custom_image_path,
                    class_names=class_names)

data/pizza_steak_sushi/04-pizza-dad.jpeg already exists, skipping
download.

```

Pred: pizza | Prob: 0.988

